

project2

December 1, 2023

1 PROJECT 2

1.1 Customer Segmentation using RFM Analysis

1.1.1 I. DATA PREPROCESSING

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
```

IMPORTING THE DATASET

```
[2]: data = pd.read_csv('data.csv')
```

```
[3]: #Displaying a few rows of data-
data.head()
```

```
[3]: InvoiceNo StockCode Description Quantity \
0 536365 85123A WHITE HANGING HEART T-LIGHT HOLDER 6
1 536365 71053 WHITE METAL LANTERN 6
2 536365 84406B CREAM CUPID HEARTS COAT HANGER 8
3 536365 84029G KNITTED UNION FLAG HOT WATER BOTTLE 6
4 536365 84029E RED WOOLLY HOTTIE WHITE HEART. 6
```

```
InvoiceDate UnitPrice CustomerID Country
0 12/1/10 8:26 2.55 17850.0 United Kingdom
1 12/1/10 8:26 3.39 17850.0 United Kingdom
2 12/1/10 8:26 2.75 17850.0 United Kingdom
3 12/1/10 8:26 3.39 17850.0 United Kingdom
4 12/1/10 8:26 3.39 17850.0 United Kingdom
```

DATA CLEANING

```
[4]: data.isnull().sum()
```

```
[4]: InvoiceNo      0
      StockCode     0
      Description   1454
      Quantity      0
      InvoiceDate    0
      UnitPrice     0
      CustomerID    135080
      Country       0
      dtype: int64
```

There are many null values in 2 columns - Description and CustomerID.

Handling the null values-

```
[5]: data['Description'].fillna('Unknown', inplace=True)
      data['Description'].isnull().sum()
```

```
[5]: 0
```

```
[6]: data['CustomerID'] = data['CustomerID'].fillna( method = 'ffill')
```

```
[7]: data.isnull().sum()
```

```
[7]: InvoiceNo      0
      StockCode     0
      Description   0
      Quantity      0
      InvoiceDate    0
      UnitPrice     0
      CustomerID    0
      Country       0
      dtype: int64
```

Null values are handled.

FORMATTING

checking if any data types needs to be converted-

```
[8]: data[['Date', 'Time']] = data['InvoiceDate'].str.split(expand=True)
```

```
[9]: data['Time']
```

```
[9]: 0      8:26
      1      8:26
      2      8:26
      3      8:26
      4      8:26
```

...

```
541904    12:50
541905    12:50
541906    12:50
541907    12:50
541908    12:50
Name: Time, Length: 541909, dtype: object
```

No changes needed here.

```
[10]: data['Date']
```

```
[10]: 0      12/1/10
      1      12/1/10
      2      12/1/10
      3      12/1/10
      4      12/1/10
      ...
541904    12/9/11
541905    12/9/11
541906    12/9/11
541907    12/9/11
541908    12/9/11
Name: Date, Length: 541909, dtype: object
```

Changing to YYYY-MM-DD format-

```
[11]: data['Date'] = pd.to_datetime(data['Date'])
      data['Date']
```

```
[11]: 0      2010-12-01
      1      2010-12-01
      2      2010-12-01
      3      2010-12-01
      4      2010-12-01
      ...
541904    2011-12-09
541905    2011-12-09
541906    2011-12-09
541907    2011-12-09
541908    2011-12-09
Name: Date, Length: 541909, dtype: datetime64[ns]
```

Data Preprocessing is completed.

1.1.2 II. RFM CALCULATION

Calculating Recency -

```
[12]: max_date = data['Date'].max()
data['Recency'] = (max_date - data['Date']).dt.days
```

```
[13]: data['Recency']
```

```
[13]: 0          373
      1          373
      2          373
      3          373
      4          373
      ...
      541904      0
      541905      0
      541906      0
      541907      0
      541908      0
      Name: Recency, Length: 541909, dtype: int64
```

Calculating Frequency -

```
[14]: frequency = data.groupby('CustomerID')['InvoiceNo'].nunique().reset_index()
frequency.columns = ['CustomerID', 'Frequency']
```

```
[15]: frequency
```

```
[15]:   CustomerID  Frequency
0      12346.0           2
1      12347.0           7
2      12348.0           5
3      12349.0           1
4      12350.0           1
...         ...         ...
4367     18280.0           6
4368     18281.0           1
4369     18282.0           3
4370     18283.0          16
4371     18287.0           3
```

[4372 rows x 2 columns]

Calculating Monetary-

```
[16]: data['TotalPrice'] = data['Quantity'] * data['UnitPrice'] # Calculate total price per transaction
monetary = data.groupby('CustomerID')['TotalPrice'].sum().reset_index()
monetary.columns = ['CustomerID', 'Monetary']
```

```
[17]: monetary
```

```
[17]:
```

| | CustomerID | Monetary |
|------|------------|----------|
| 0 | 12346.0 | 0.00 |
| 1 | 12347.0 | 4310.00 |
| 2 | 12348.0 | 3366.27 |
| 3 | 12349.0 | 1757.55 |
| 4 | 12350.0 | 334.40 |
| ... | ... | ... |
| 4367 | 18280.0 | 8330.57 |
| 4368 | 18281.0 | 80.82 |
| 4369 | 18282.0 | 176.60 |
| 4370 | 18283.0 | 2094.88 |
| 4371 | 18287.0 | 1837.28 |

[4372 rows x 2 columns]

```
[18]: rfm = pd.merge(frequency, monetary, on='CustomerID')
rfm = pd.merge(rfm, data[['CustomerID', 'Recency']].drop_duplicates(),
               ↪on='CustomerID')
```

RFM data frame-

```
[19]: rfm
```

```
[19]:
```

| | CustomerID | Frequency | Monetary | Recency |
|-------|------------|-----------|----------|---------|
| 0 | 12346.0 | 2 | 0.00 | 325 |
| 1 | 12347.0 | 7 | 4310.00 | 367 |
| 2 | 12347.0 | 7 | 4310.00 | 317 |
| 3 | 12347.0 | 7 | 4310.00 | 246 |
| 4 | 12347.0 | 7 | 4310.00 | 183 |
| ... | ... | ... | ... | ... |
| 19300 | 18283.0 | 16 | 2094.88 | 9 |
| 19301 | 18283.0 | 16 | 2094.88 | 3 |
| 19302 | 18287.0 | 3 | 1837.28 | 201 |
| 19303 | 18287.0 | 3 | 1837.28 | 58 |
| 19304 | 18287.0 | 3 | 1837.28 | 42 |

[19305 rows x 4 columns]

1.1.3 III. RFM SEGMENTATION

Using quartiles-

```
[20]: rfm['RecencyScore'] = pd.qcut(rfm['Recency'], q=4, labels=range(4, 0, -1))
rfm['FrequencyScore'] = pd.qcut(rfm['Frequency'], q=4, labels=range(1, 5))
rfm['MonetaryScore'] = pd.qcut(rfm['Monetary'], q=4, labels=range(1, 5))
```

```
[21]: rfm['RecencyScore'] = rfm['RecencyScore'].astype(int)
rfm['FrequencyScore'] = rfm['FrequencyScore'].astype(int)
```

```
rfm['MonetaryScore'] = rfm['MonetaryScore'].astype(int)
```

Calculating the score-

```
[22]: rfm['RFMScore'] = rfm['RecencyScore'] + rfm['FrequencyScore'] +
      ↪rfm['MonetaryScore']
```

```
[23]: rfm
```

```
[23]:
```

| | CustomerID | Frequency | Monetary | Recency | RecencyScore | FrequencyScore | \ |
|-------|------------|-----------|----------|---------|--------------|----------------|---|
| 0 | 12346.0 | 2 | 0.00 | 325 | 1 | 1 | |
| 1 | 12347.0 | 7 | 4310.00 | 367 | 1 | 2 | |
| 2 | 12347.0 | 7 | 4310.00 | 317 | 1 | 2 | |
| 3 | 12347.0 | 7 | 4310.00 | 246 | 2 | 2 | |
| 4 | 12347.0 | 7 | 4310.00 | 183 | 2 | 2 | |
| ... | ... | ... | ... | ... | ... | ... | |
| 19300 | 18283.0 | 16 | 2094.88 | 9 | 4 | 3 | |
| 19301 | 18283.0 | 16 | 2094.88 | 3 | 4 | 3 | |
| 19302 | 18287.0 | 3 | 1837.28 | 201 | 2 | 1 | |
| 19303 | 18287.0 | 3 | 1837.28 | 58 | 4 | 1 | |
| 19304 | 18287.0 | 3 | 1837.28 | 42 | 4 | 1 | |

| | MonetaryScore | RFMScore |
|-------|---------------|----------|
| 0 | 1 | 3 |
| 1 | 3 | 6 |
| 2 | 3 | 6 |
| 3 | 3 | 7 |
| 4 | 3 | 7 |
| ... | ... | ... |
| 19300 | 2 | 9 |
| 19301 | 2 | 9 |
| 19302 | 2 | 5 |
| 19303 | 2 | 7 |
| 19304 | 2 | 7 |

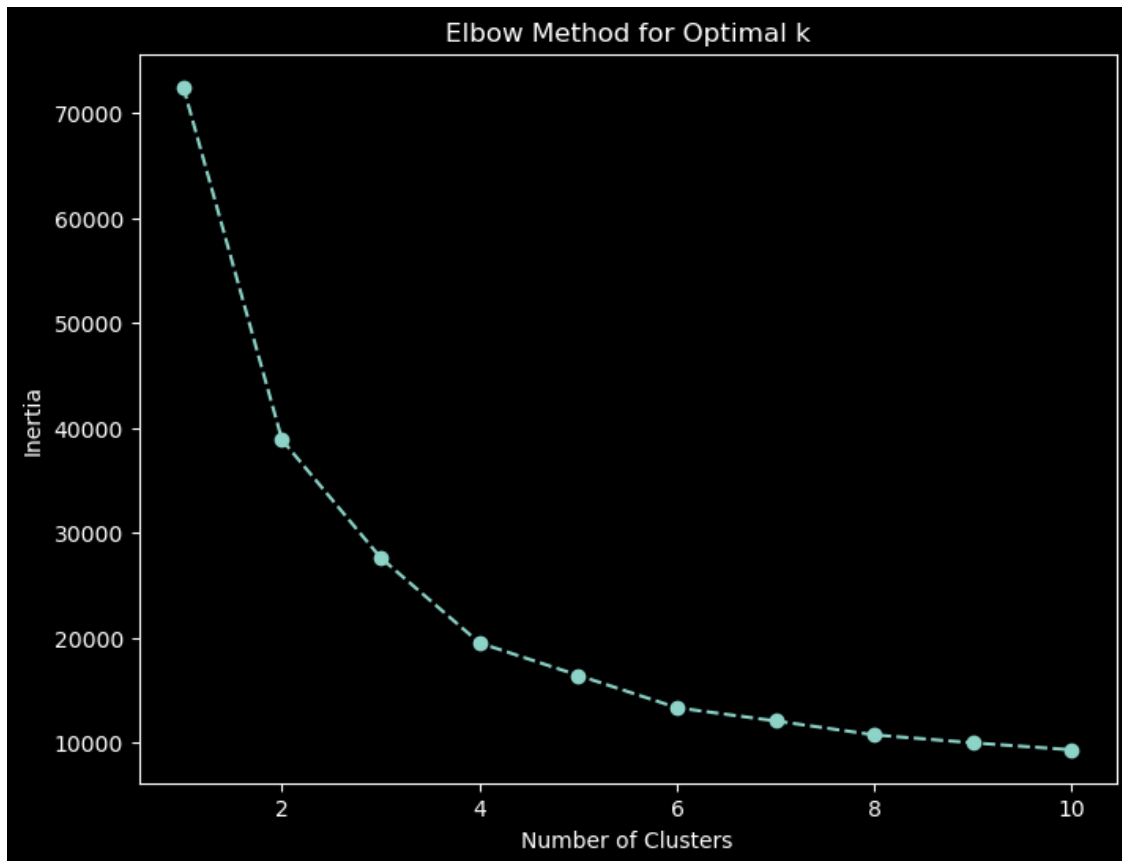
[19305 rows x 8 columns]

1.1.4 IV. CUSTOMER SEGMENTATION

```
[24]: X = rfm[['RecencyScore', 'FrequencyScore', 'MonetaryScore']]
```

```
[25]: inertia = []
      for k in range(1, 11):
          kmeans = KMeans(n_clusters=k, n_init=10, random_state=42)
          kmeans.fit(X)
          inertia.append(kmeans.inertia_)
```

```
plt.style.use('dark_background')
plt.figure(figsize=(8, 6))
plt.plot(range(1, 11), inertia, marker='o', linestyle='--')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.title('Elbow Method for Optimal k')
plt.show()
```



From the Elbow method, the optimal number of clusters is 6

```
[26]: kmeans = KMeans(n_clusters=6, n_init=10, random_state=42)
kmeans.fit(X)
rfm['Cluster'] = kmeans.labels_
print(rfm.head())
```

| | CustomerID | Frequency | Monetary | Recency | RecencyScore | FrequencyScore | \ |
|---|------------|-----------|----------|---------|--------------|----------------|---|
| 0 | 12346.0 | 2 | 0.0 | 325 | 1 | 1 | |
| 1 | 12347.0 | 7 | 4310.0 | 367 | 1 | 2 | |
| 2 | 12347.0 | 7 | 4310.0 | 317 | 1 | 2 | |
| 3 | 12347.0 | 7 | 4310.0 | 246 | 2 | 2 | |
| 4 | 12347.0 | 7 | 4310.0 | 183 | 2 | 2 | |

| | MonetaryScore | RFMScore | Cluster |
|---|---------------|----------|---------|
| 0 | 1 | 3 | 4 |
| 1 | 3 | 6 | 2 |
| 2 | 3 | 6 | 2 |
| 3 | 3 | 7 | 2 |
| 4 | 3 | 7 | 2 |

```
[27]: rfm
```

```
[27]:
```

| | CustomerID | Frequency | Monetary | Recency | RecencyScore | FrequencyScore | \ |
|-------|------------|-----------|----------|---------|--------------|----------------|---|
| 0 | 12346.0 | 2 | 0.00 | 325 | 1 | 1 | |
| 1 | 12347.0 | 7 | 4310.00 | 367 | 1 | 2 | |
| 2 | 12347.0 | 7 | 4310.00 | 317 | 1 | 2 | |
| 3 | 12347.0 | 7 | 4310.00 | 246 | 2 | 2 | |
| 4 | 12347.0 | 7 | 4310.00 | 183 | 2 | 2 | |
| ... | ... | ... | ... | ... | ... | ... | |
| 19300 | 18283.0 | 16 | 2094.88 | 9 | 4 | 3 | |
| 19301 | 18283.0 | 16 | 2094.88 | 3 | 4 | 3 | |
| 19302 | 18287.0 | 3 | 1837.28 | 201 | 2 | 1 | |
| 19303 | 18287.0 | 3 | 1837.28 | 58 | 4 | 1 | |
| 19304 | 18287.0 | 3 | 1837.28 | 42 | 4 | 1 | |

| | MonetaryScore | RFMScore | Cluster |
|-------|---------------|----------|---------|
| 0 | 1 | 3 | 4 |
| 1 | 3 | 6 | 2 |
| 2 | 3 | 6 | 2 |
| 3 | 3 | 7 | 2 |
| 4 | 3 | 7 | 2 |
| ... | ... | ... | ... |
| 19300 | 2 | 9 | 5 |
| 19301 | 2 | 9 | 5 |
| 19302 | 2 | 5 | 4 |
| 19303 | 2 | 7 | 1 |
| 19304 | 2 | 7 | 1 |

```
[19305 rows x 9 columns]
```

1.1.5 V. SEGMENT PROFILING

```
[28]: segment_profiles = rfm.groupby('Cluster').agg({
    'Recency': 'mean',
    'Frequency': 'mean',
    'Monetary': 'mean',
}).reset_index()

segment_profiles.rename(columns={'Cluster': 'Segment'}, inplace=True)
```



```
print(segment_profiles)
```

| | Segment | Recency | Frequency | Monetary |
|---|---------|------------|-----------|--------------|
| 0 | 0 | 262.507607 | 57.863071 | 27742.172801 |
| 1 | 1 | 60.830386 | 2.984283 | 608.319043 |
| 2 | 2 | 258.112774 | 11.153443 | 3031.899866 |
| 3 | 3 | 68.897812 | 43.802317 | 21316.296412 |
| 4 | 4 | 262.752304 | 3.177295 | 530.321324 |
| 5 | 5 | 65.426444 | 8.369778 | 2332.801885 |

SEGMENT PROFILING: RFM Analysis:

1. Recency: Mean recency values for each segment indicate how recently customers made purchases. Lower values signify more recent purchases.
2. Frequency (F): Mean or median frequency values demonstrate how often customers from each segment make purchases. Higher values indicate more frequent buyers.
3. Monetary (M): Mean or median monetary values represent the average spending of customers within each segment. Higher values indicate higher-spending customers.

Segment 0:

- Recency (R): Around 69 days since last purchase.
- Frequency (F): High frequency, averaging 44 orders.
- Monetary (M): High monetary value, spending about \$21,287 on average.
- Profile: Engaged and high-spending customers who make frequent purchases.

Segment 1:

- Recency (R): Approximately 61 days since last purchase.
- Frequency (F): Low frequency, averaging around 3 orders.
- Monetary (M): Lower monetary value, spending about \$608 on average.
- Profile: Customers who make fewer purchases, less engaged, and with lower spending.

Segment 2:

- Recency (R): Higher recency, about 258 days since last purchase.
- Frequency (F): Moderate frequency, averaging around 11 orders.
- Monetary (M): Moderate monetary value, spending about \$3,033 on average.
- Profile: Customers with moderate engagement and spending, but less recent activity.

Segment 3:

- Recency (R): Approximately 65 days since last purchase.
- Frequency (F): Moderate frequency, averaging about 8 orders.
- Monetary (M): Moderate monetary value, spending around \$2,333 on average.
- Profile: Moderately engaged customers with moderate spending patterns.

Segment 4:

- Recency (R): High recency, around 263 days since last purchase.
- Frequency (F): Low frequency, averaging about 3 orders.
- Monetary (M): Lower monetary value, spending approximately \$530 on average.
- Profile: Customers who make infrequent purchases and exhibit low spending habits.

Segment 5:

- Recency (R): Higher recency, about 262 days since last purchase.
- Frequency (F): High frequency, averaging 58 orders.
- Monetary (M): High monetary value, spending approximately \$27,747 on average.
- Profile: Engaged and high-spending customers with frequent purchase behavior.

1.1.6 VI. MARKETING RECOMMENDATIONS**Segment 0: Engaged High-Spending Customers****Recommendations:**

- VIP Treatment: Offer exclusive perks, early access, or VIP rewards to maintain loyalty.
- Personalized Offers: Create personalized bundles, discounts, or promotions based on their preferences and purchase history.
- Loyalty Programs: Implement tiered loyalty programs to encourage repeat purchases and increase engagement.
- Upselling/Cross-selling: Suggest complementary high-value products to increase the average order value.

Segment 1: Low-Engagement, Low-Spending Customers**Recommendations:**

- Re-engagement Campaigns: Send targeted campaigns with incentives or discounts to encourage revisits.
- Personalized Recommendations: Recommend products based on past purchases to spark interest.
- Improve Customer Experience: Enhance user experience, ease of ordering, or customer service to improve satisfaction and retention.

Segment 2: Moderately Engaged, Moderate Spending**Recommendations:**

- Reactivation Campaigns: Target customers with personalized incentives to re-engage.
- Showcase Value: Highlight unique features, benefits, or promotions to encourage more frequent purchases.
- Loyalty Incentives: Offer loyalty rewards or discounts on subsequent purchases.

Segment 3: Moderate Engagement, Moderate Spending

Recommendations:

- Tailored Promotions: Send personalized promotions based on past behavior to encourage increased spending.
- Loyalty Rewards: Encourage loyalty with rewards or points for continued engagement.
- Product Education: Share educational content or tips related to their purchases to enhance engagement.

Segment 4: Low-Engagement, Low-Spending

Recommendations:

- Win-Back Campaigns: Target with compelling offers to regain their interest.
- Incentivize Purchases: Offer discounts or incentives to increase order frequency.
- Referral Programs: Encourage referrals by offering benefits for recommending products to others.

Segment 5: Engaged High-Frequency, High-Spending Customers


Recommendations:

- Exclusive Access: Provide early access to new products, services, or events to maintain engagement.
- Reward Loyalty: Offer special rewards or benefits for continued loyalty and high spending.
- Personalized Engagement: Tailor communications and offers based on past interactions to enhance satisfaction and retention.


1.1.7 VISUALIZATION

```
[29]: import plotly.express as px


fig = px.bar(segment_profiles, x='Segment', y='Recency', color='Segment',
             title='Recency by Segment')
fig.update_layout( width=800, height=600,
                  plot_bgcolor='black',
                  paper_bgcolor='black',
                  font=dict(color='white')
                )
fig.show()
```




```
[30]: fig = px.bar(segment_profiles, x='Segment', y='Frequency', color='Segment',  
    ↪title='Frequency by Segment')  
fig.update_layout( width=800, height=600,  
    plot_bgcolor='black',  
    paper_bgcolor='black',  
    font=dict(color='white')  
)  
fig.show()
```



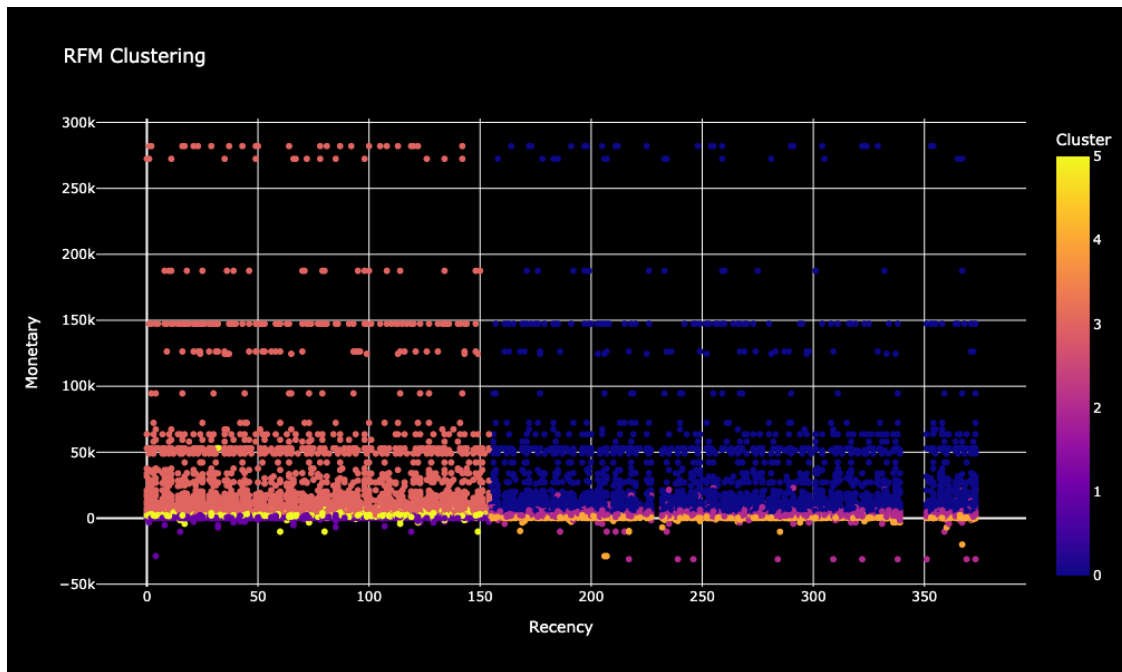
```
[31]: fig = px.bar(segment_profiles, x='Segment', y='Monetary', color='Segment', title='Monetary by Segment')
fig.update_layout( width=800, height=600,
    plot_bgcolor='black',
    paper_bgcolor='black',
    font=dict(color='white')
)
fig.show()
```



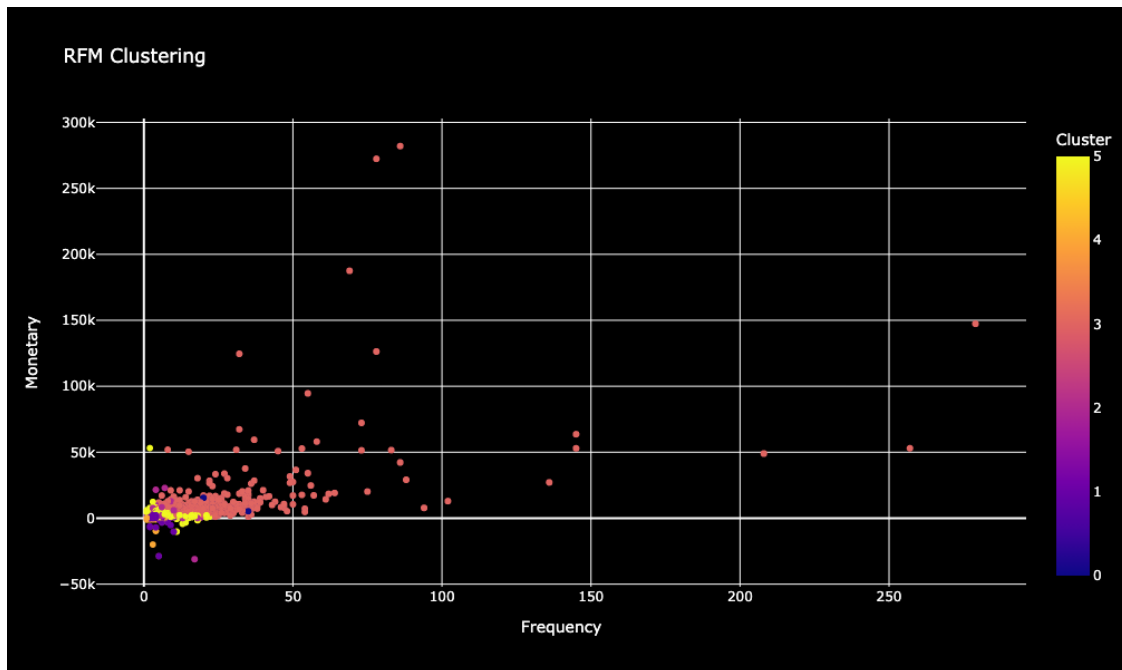
```
[32]: fig = px.scatter(rfm, x='Recency', y='Frequency' , color='Cluster', title='RFM_
      ↳Clustering')
fig.update_layout( width=800, height=600,
                    plot_bgcolor='black',
                    paper_bgcolor='black',
                    font=dict(color='white')
                )
fig.show()
```



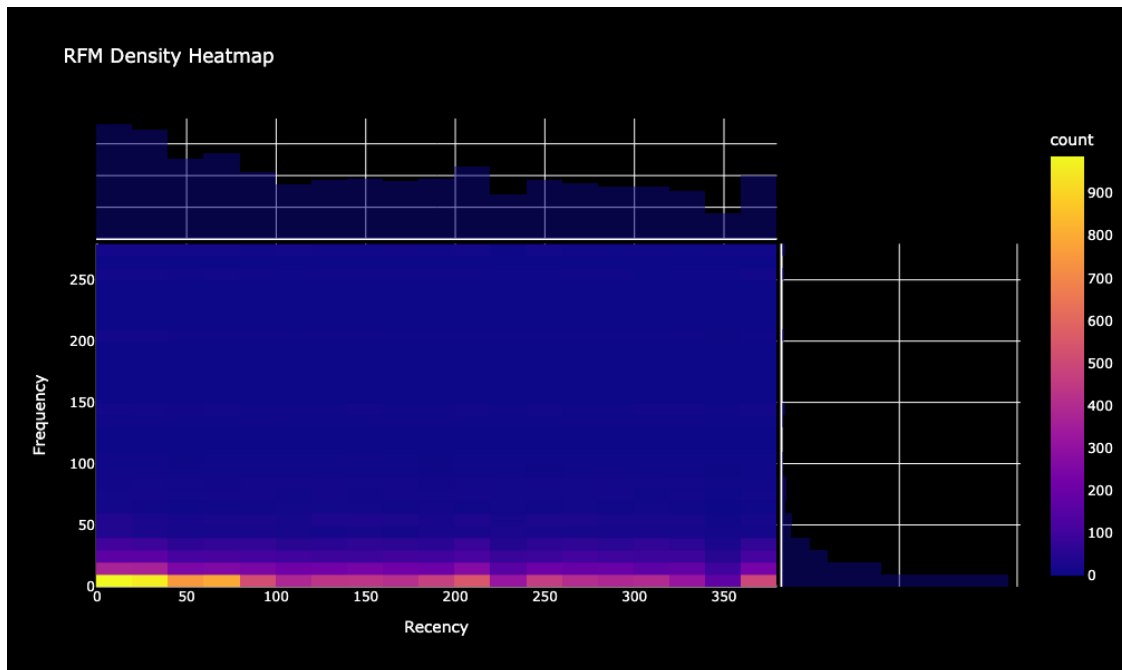
```
[33]: fig = px.scatter(rfm, x='Recency', y='Monetary' , color='Cluster', title='RFM_
      ↳Clustering')
fig.update_layout( width=800, height=600,
                    plot_bgcolor='black',
                    paper_bgcolor='black',
                    font=dict(color='white')
                )
fig.show()
```




```
[34]: fig = px.scatter(rfm, x='Frequency', y='Monetary' , color='Cluster', title='RFM Clustering')
fig.update_layout( width=800, height=600,
    plot_bgcolor='black',
    paper_bgcolor='black',
    font=dict(color='white')
)
fig.show()
```


```
[35]: fig = px.density_heatmap(rfm, x='Recency', y='Frequency',
    ↪marginal_x='histogram', marginal_y='histogram', title='RFM Density Heatmap')
fig.update_layout( width=800, height=600,
    plot_bgcolor='black',
    paper_bgcolor='black',
    font=dict(color='white')
)
fig.show()
```



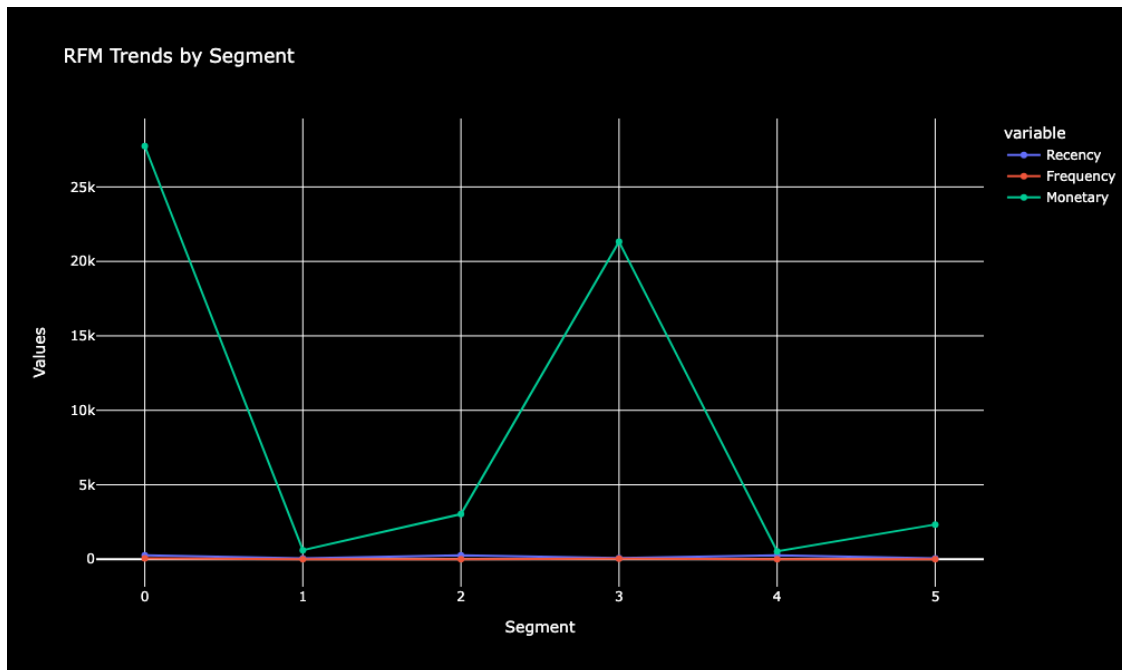
```
[36]: fig = px.density_heatmap(rfm, x='Recency', y='Monetary',
    ↪marginal_x='histogram', marginal_y='histogram', title='RFM Density Heatmap')
fig.update_layout( width=800, height=600,
    plot_bgcolor='black',
    paper_bgcolor='black',
    font=dict(color='white')
)
fig.show()
```




```
[37]: fig = px.density_heatmap(rfm, x='Monetary', y='Frequency',
    ↪marginal_x='histogram', marginal_y='histogram', title='RFM Density Heatmap')
fig.update_layout(width=800, height=600,
    plot_bgcolor='black',
    paper_bgcolor='black',
    font=dict(color='white')
)
fig.show()
```



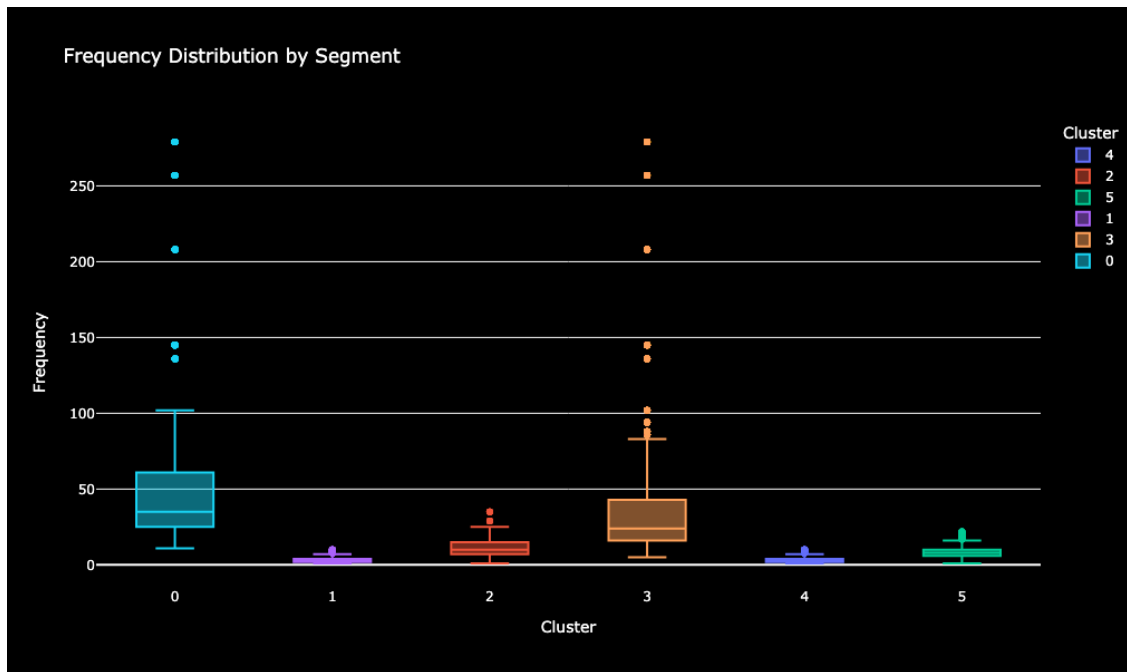
```
[38]: fig = px.line(segment_profiles, x='Segment', y=['Recency', 'Frequency', 'Monetary'], title='RFM Trends by Segment')
fig.update_traces(mode='lines+markers')
fig.update_layout( width=800, height=600,
    plot_bgcolor='black',
    paper_bgcolor='black',
    font=dict(color='white'),
    xaxis=dict(type='category'),
    yaxis=dict(title='Values')
)
fig.show()
```




```
[39]: fig = px.box(rfm, x='Cluster', y='Monetary', color='Cluster', title='Monetary_
      ↳ Distribution by Segment')
fig.update_layout( width=800, height=600,
                  plot_bgcolor='black',
                  paper_bgcolor='black',
                  font=dict(color='white')
                )
fig.show()
```




```
[40]: fig = px.box(rfm, x='Cluster', y='Frequency', color='Cluster', title='Frequency_
      ↳Distribution by Segment')
fig.update_layout( width=800, height=600,
                    plot_bgcolor='black',
                    paper_bgcolor='black',
                    font=dict(color='white')
                )
fig.show()
```



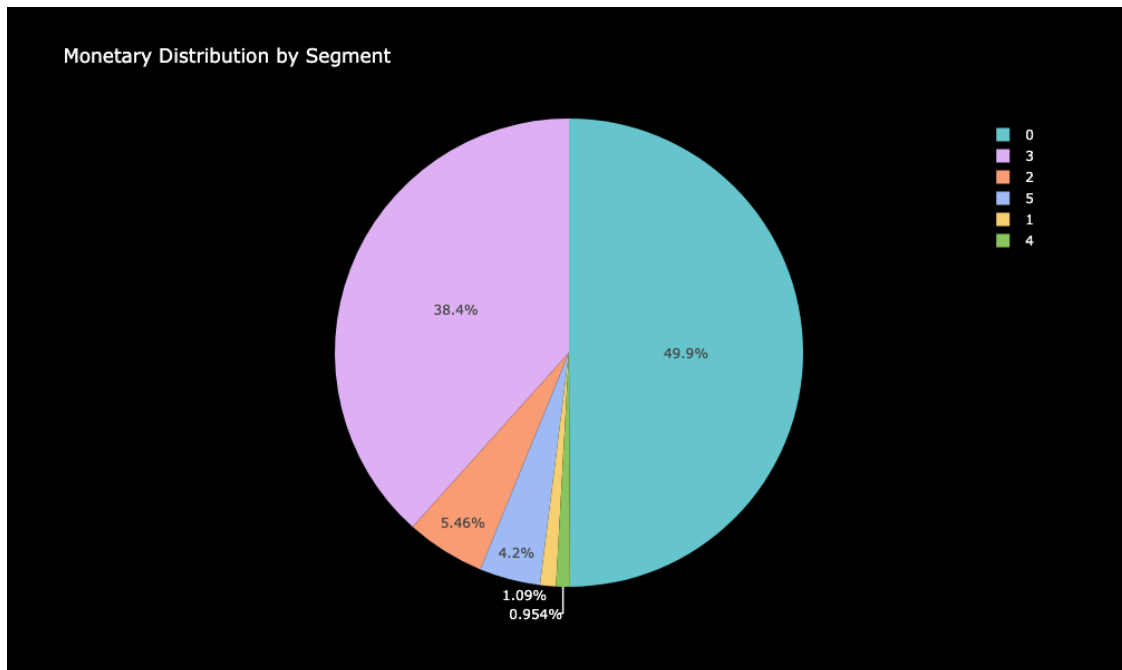
```
[41]: fig = px.box(rfm, x='Cluster', y='Recency', color='Cluster', title='Recency_
      ↪Distribution by Segment')
fig.update_layout( width=800, height=600,
      plot_bgcolor='black',
      paper_bgcolor='black',
      font=dict(color='white')
)
fig.show()
```




```
[42]: fig_recency = px.pie(segment_profiles, names='Segment', values='Recency',
    ↪title='Recency Distribution by Segment')
fig_recency.update_traces(marker=dict(colors=px.colors.qualitative.Pastel))
fig_recency.update_layout( width=800, height=600,
    plot_bgcolor='black',
    paper_bgcolor='black',
    font=dict(color='white')
)
fig_recency.show()
```


```
[43]: fig_monetary = px.pie(segment_profiles, names='Segment', values='Monetary',
    ↪title='Monetary Distribution by Segment')
fig_monetary.update_traces(marker=dict(colors=px.colors.qualitative.Pastel))
fig_monetary.update_layout( width=800, height=600,
    plot_bgcolor='black',
    paper_bgcolor='black',
    font=dict(color='white')
)
fig_monetary.show()
```



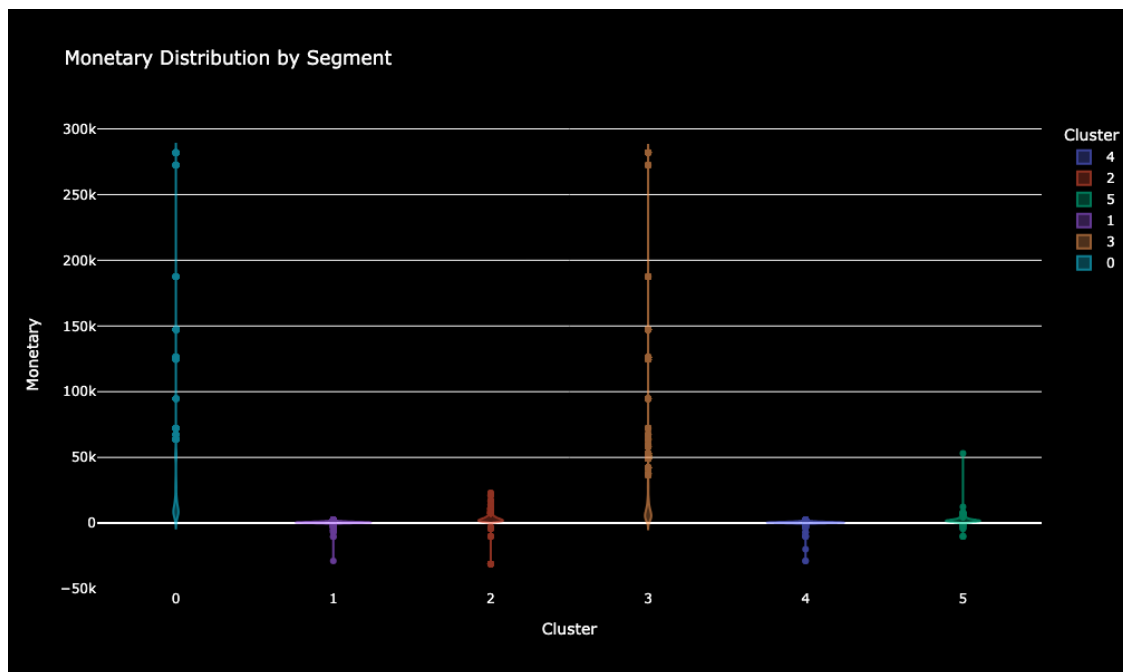
```
[44]: fig_frequency = px.pie(segment_profiles, names='Segment', values='Frequency',
    ↪title='Frequency Distribution by Segment')
fig_frequency.update_traces(marker=dict(colors=px.colors.qualitative.Pastel))
fig_frequency.update_layout( width=800, height=600,
    plot_bgcolor='black',
    paper_bgcolor='black',
    font=dict(color='white')
)
fig_frequency.show()
```



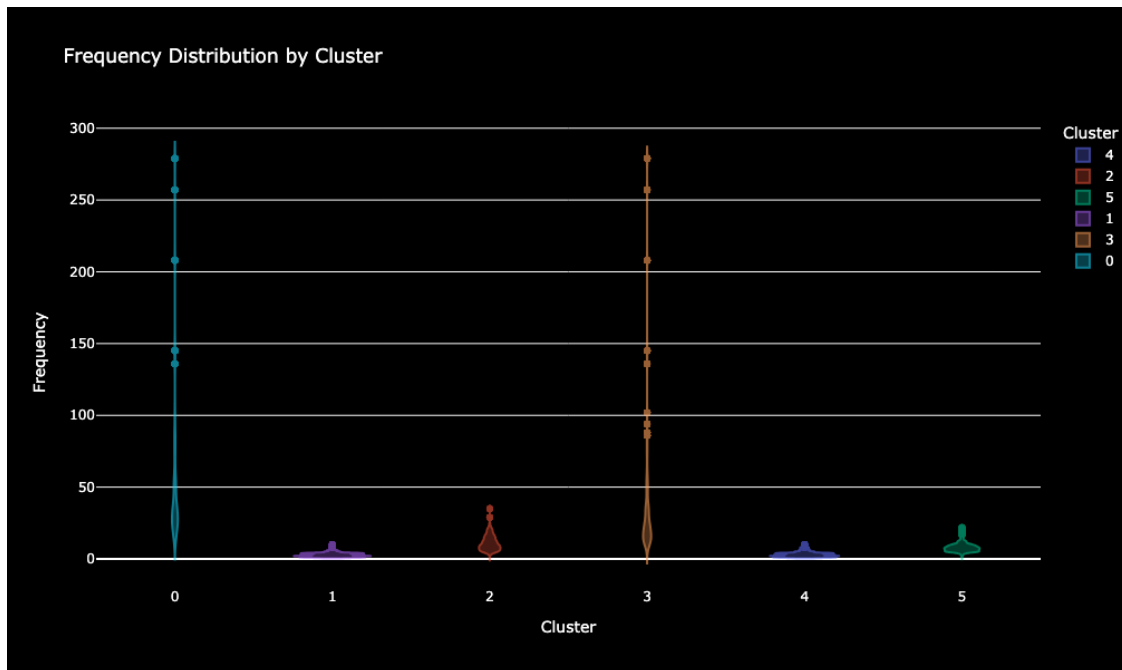
```
[45]: fig_recency_violin = px.violin(rfm, x='Cluster', y='Recency', color='Cluster',
    ↪title='Recency Distribution by Segment')
fig_recency_violin.update_traces(marker=dict(line=dict(color='black')),
    ↪opacity=0.6)
fig_recency_violin.update_layout( width=800, height=600,
    plot_bgcolor='black',
    paper_bgcolor='black',
    font=dict(color='white')
)
fig_recency_violin.show()
```



```
[46]: fig_recency_violin = px.violin(rfm, x='Cluster', y='Monetary', color='Cluster',
    ↪title='Monetary Distribution by Segment')
fig_recency_violin.update_traces(marker=dict(line=dict(color='black')),
    ↪opacity=0.6)
fig_recency_violin.update_layout( width=800, height=600,
    plot_bgcolor='black',
    paper_bgcolor='black',
    font=dict(color='white')
)
fig_recency_violin.show()
```




```
[47]: fig_recency_violin = px.violin(rfm, x='Cluster', y='Frequency',
    ↪color='Cluster', title='Frequency Distribution by Cluster')
fig_recency_violin.update_traces(marker=dict(line=dict(color='black')),
    ↪opacity=0.6)
fig_recency_violin.update_layout( width=800, height=600,
    plot_bgcolor='black',
    paper_bgcolor='black',
    font=dict(color='white')
)
fig_recency_violin.show()
```



```
[48]: import plotly.graph_objects as go
fig = go.Figure(data=[go.Scatter3d(
    x=rfm['Recency'],
    y=rfm['Frequency'],
    z=rfm['Monetary'],
    mode='markers',
    marker=dict(
        size=8,
        color=rfm['Recency'],
        colorscale='Viridis',
        colorbar=dict(title='Recency'),
        line=dict(color='black', width=0.5)
    )
)])

fig.update_layout(
    scene=dict(
        xaxis=dict(title='Recency'),
        yaxis=dict(title='Frequency'),
        zaxis=dict(title='Monetary'),
    ),
    margin=dict(l=0, r=0, b=0, t=0),
)

fig.show()
```

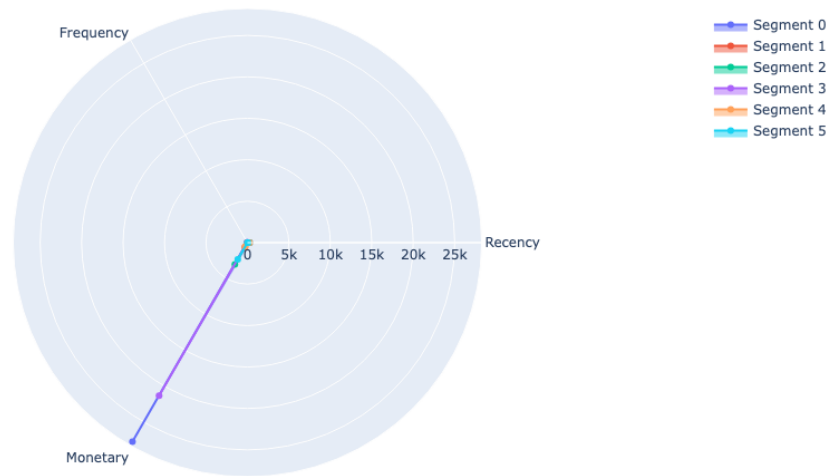


```
[49]: fig = go.Figure()

for segment in segment_profiles['Segment']:
    fig.add_trace(go.Scatterpolar(
        r=segment_profiles[segment_profiles['Segment'] == segment][['Recency', 'Frequency', 'Monetary']].values.flatten(),
        theta=['Recency', 'Frequency', 'Monetary'],
        fill='toself',
        name=f'Segment {segment}'
    ))

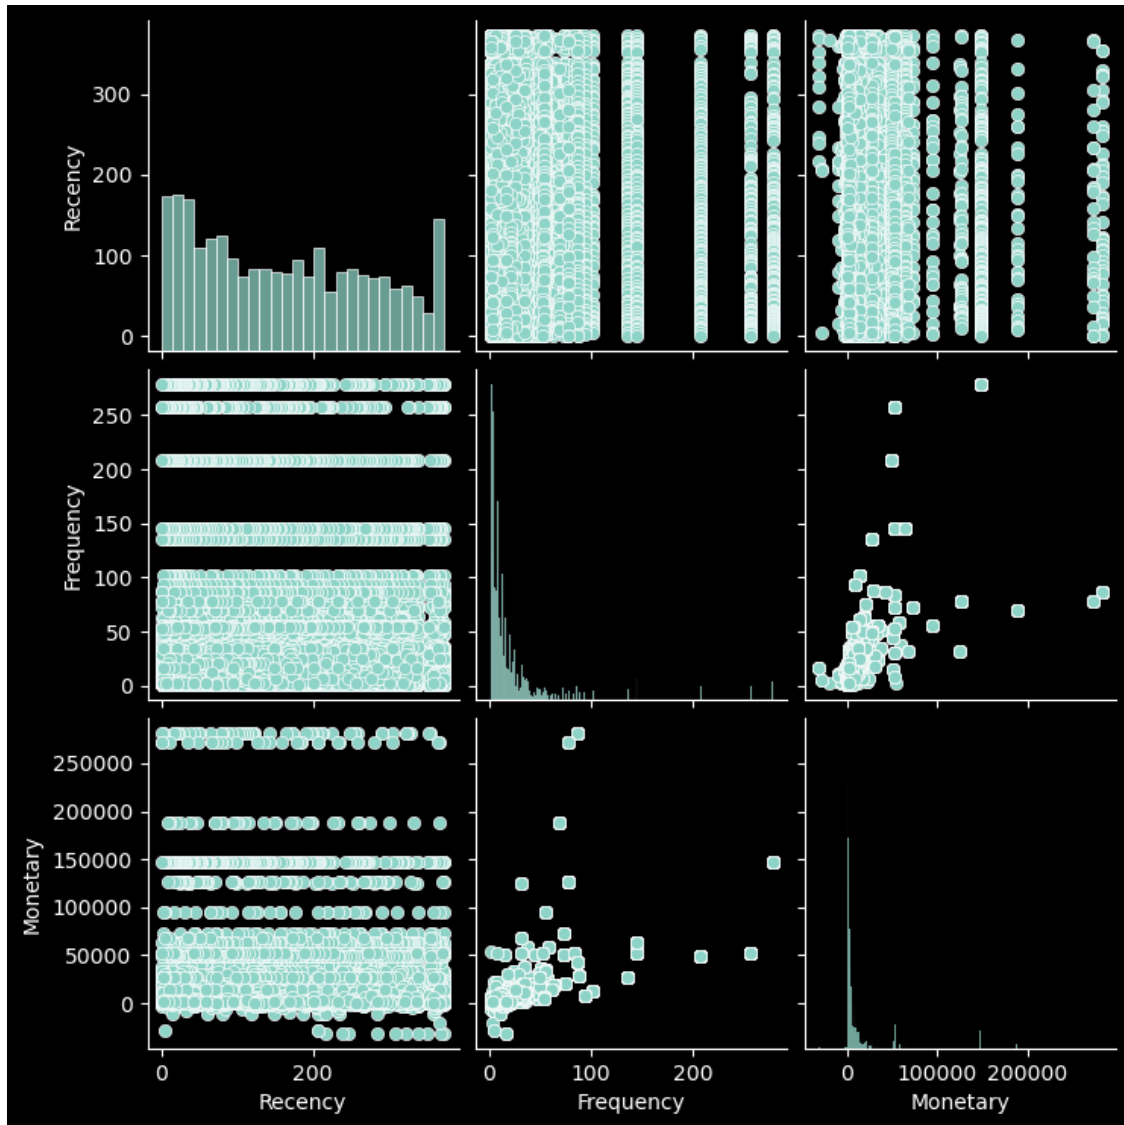
fig.update_layout( width=800, height=600,
    polar=dict(radialaxis=dict(visible=True)),
    showlegend=True
)

fig.show()
```




```
[50]: import seaborn as sns
import matplotlib.pyplot as plt

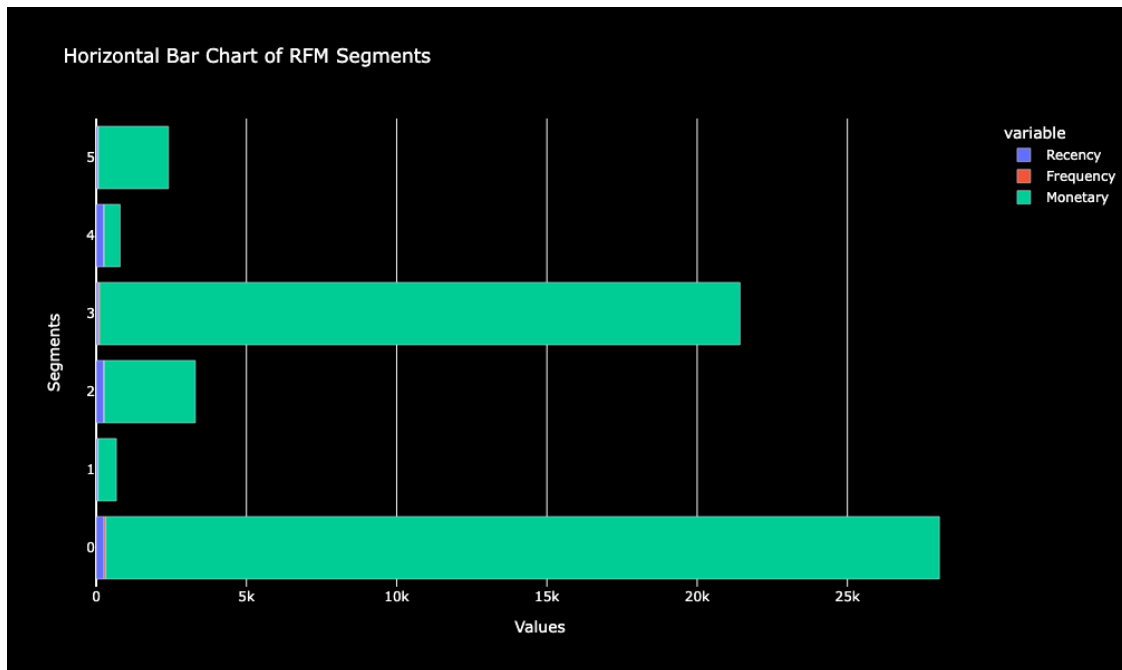
sns.pairplot(rfm[['Recency', 'Frequency', 'Monetary']])
plt.show()
```

```
[51]: sns.jointplot(x='Recency', y='Frequency', data=rfm, kind='scatter')
plt.show()
```



```
[52]: fig = px.bar(segment_profiles, y='Segment', x=['Recency', 'Frequency',
    ↳ 'Monetary'], orientation='h', title='Horizontal Bar Chart of RFM Segments')
fig.update_layout(width=800, height=600, xaxis=dict(title='Values'),
    ↳ yaxis=dict(title='Segments'), plot_bgcolor='black', paper_bgcolor='black',
    ↳ font=dict(color='white'))
fig.show()
```



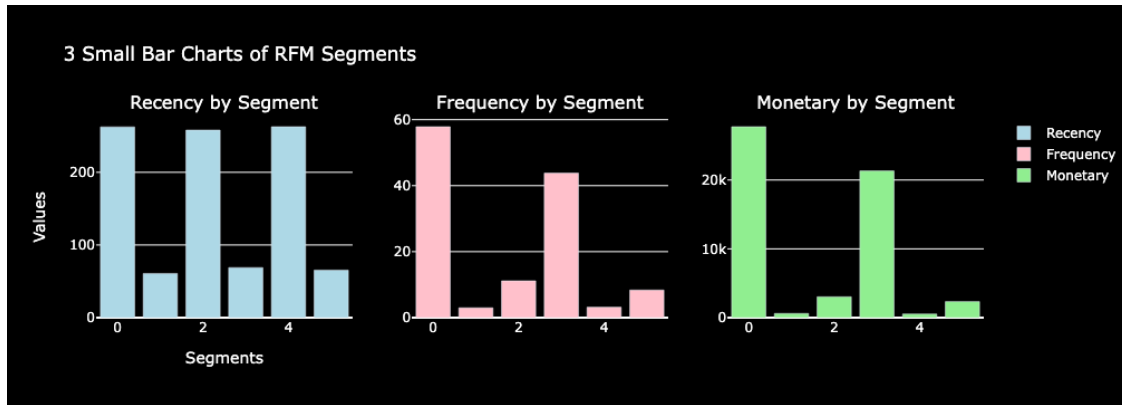
```
[53]: from plotly.subplots import make_subplots

fig = make_subplots(rows=1, cols=3, subplot_titles=('Recency by Segment',
    ↪ 'Frequency by Segment', 'Monetary by Segment'))


colors = ['lightblue', 'pink', 'lightgreen']

for i, col in enumerate(['Recency', 'Frequency', 'Monetary']):
    trace = go.Bar(x=segment_profiles['Segment'], y=segment_profiles[col],
    ↪ name=col, marker_color=colors[i])
    fig.add_trace(trace, row=1, col=i+1)

fig.update_layout(title='3 Small Bar Charts of RFM Segments',
    ↪ xaxis=dict(title='Segments'), yaxis=dict(title='Values'),
    ↪ plot_bgcolor='black', paper_bgcolor='black', font=dict(color='white'))
fig.show()
```



```
[54]: fig = px.bar(segment_profiles, x='Segment', y=['Recency', 'Frequency', 'Monetary'], barmode='stack', title='Stacked Bar Chart of RFM Segments')
fig.update_layout(xaxis=dict(title='Segments'), yaxis=dict(title='Values'))
fig.show()
```



1.1.8 SOLUTIONS-

1. DATA OVERVIEW

Q. What is the size of the dataset in terms of the number of rows and columns?

```
[55]: print(data.size)
```

6502908

We have 6502908 values in the dataset.

```
[56]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
```

Data columns (total 12 columns):

| # | Column | Non-Null Count | Dtype |
|----|-------------|-----------------|----------------|
| 0 | InvoiceNo | 541909 non-null | object |
| 1 | StockCode | 541909 non-null | object |
| 2 | Description | 541909 non-null | object |
| 3 | Quantity | 541909 non-null | int64 |
| 4 | InvoiceDate | 541909 non-null | object |
| 5 | UnitPrice | 541909 non-null | float64 |
| 6 | CustomerID | 541909 non-null | float64 |
| 7 | Country | 541909 non-null | object |
| 8 | Date | 541909 non-null | datetime64[ns] |
| 9 | Time | 541909 non-null | object |
| 10 | Recency | 541909 non-null | int64 |
| 11 | TotalPrice | 541909 non-null | float64 |

dtypes: datetime64[ns](1), float64(3), int64(2), object(6)

memory usage: 49.6+ MB

Number of columns - 12

```
[57]: print(len(data))
```

541909

Number of columns - 541909

Q. Can you provide a brief description of each column in the dataset?

```
[58]: data.describe(include='all' , datetime_is_numeric=True)
```

```
[58]:
```

| | InvoiceNo | StockCode | Description | Quantity \ |
|--------|-----------|-----------|------------------------------------|---------------|
| count | 541909 | 541909 | 541909 | 541909.000000 |
| unique | 25900 | 4070 | 4224 | NaN |
| top | 573585 | 85123A | WHITE HANGING HEART T-LIGHT HOLDER | NaN |
| freq | 1114 | 2313 | 2369 | NaN |
| mean | NaN | NaN | NaN | 9.552250 |
| min | NaN | NaN | NaN | -80995.000000 |
| 25% | NaN | NaN | NaN | 1.000000 |
| 50% | NaN | NaN | NaN | 3.000000 |
| 75% | NaN | NaN | NaN | 10.000000 |
| max | NaN | NaN | NaN | 80995.000000 |
| std | NaN | NaN | NaN | 218.081158 |

| | InvoiceDate | UnitPrice | CustomerID | Country \ |
|--------|----------------|---------------|---------------|----------------|
| count | 541909 | 541909.000000 | 541909.000000 | 541909 |
| unique | 23260 | NaN | NaN | 38 |
| top | 10/31/11 14:41 | NaN | NaN | United Kingdom |
| freq | 1114 | NaN | NaN | 495478 |
| mean | NaN | 4.611114 | 15272.795237 | NaN |

| | | | | |
|-----|-----|---------------|--------------|-----|
| min | NaN | -11062.060000 | 12346.000000 | NaN |
| 25% | NaN | 1.250000 | 13798.000000 | NaN |
| 50% | NaN | 2.080000 | 15145.000000 | NaN |
| 75% | NaN | 4.130000 | 16803.000000 | NaN |
| max | NaN | 38970.000000 | 18287.000000 | NaN |
| std | NaN | 96.759853 | 1737.934523 | NaN |

| | Date | Time | Recency | TotalPrice |
|--------|-------------------------------|--------|---------------|----------------|
| count | 541909 | 541909 | 541909.000000 | 541909.000000 |
| unique | NaN | 774 | NaN | NaN |
| top | NaN | 15:56 | NaN | NaN |
| freq | NaN | 2628 | NaN | NaN |
| mean | 2011-07-04 00:00:13.073782272 | NaN | 157.999849 | 17.987795 |
| min | 2010-12-01 00:00:00 | NaN | 0.000000 | -168469.600000 |
| 25% | 2011-03-28 00:00:00 | NaN | 51.000000 | 3.400000 |
| 50% | 2011-07-19 00:00:00 | NaN | 143.000000 | 9.750000 |
| 75% | 2011-10-19 00:00:00 | NaN | 256.000000 | 17.400000 |
| max | 2011-12-09 00:00:00 | NaN | 373.000000 | 168469.600000 |
| std | NaN | NaN | 115.877074 | 378.810824 |

Q. What is the time period covered by this dataset?

```
[59]: df_search=data[data['Date'].notnull()]
start = df_search['Date'].min()
end = df_search['Date'].max()
print(f"start date: {start} ")
print(f"end date: {end}")
```

```
start date: 2010-12-01 00:00:00
end date: 2011-12-09 00:00:00
```

```
[60]: time_period = end - start
```

```
[61]: print("The time period covered in this dataset is- \n", time_period)
```

```
The time period covered in this dataset is-
373 days 00:00:00
```

2. CUSTOMER ANALYSIS

Q. How many unique customers are there in the dataset?

```
[62]: data.groupby('CustomerID')['InvoiceNo'].nunique()
```

```
[62]: CustomerID
12346.0    2
12347.0    7
12348.0    5
```

```
12349.0    1
12350.0    1
..
18280.0    6
18281.0    1
18282.0    3
18283.0   16
18287.0    3
Name: InvoiceNo, Length: 4372, dtype: int64
```

```
[63]: data['CustomerID'].nunique()
```

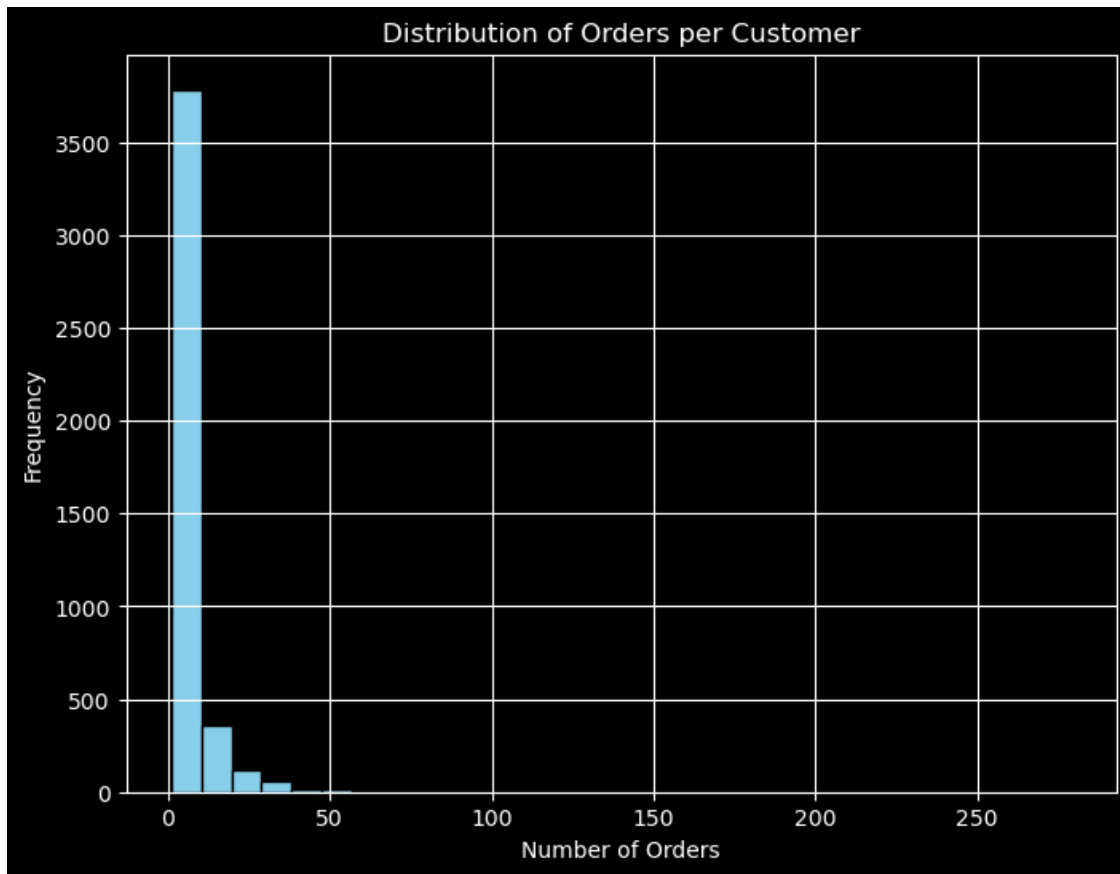
```
[63]: 4372
```

There are 4372 unique customers.

Q. What is the distribution of the number of orders per customer?

```
[64]: orders_per_customer = data.groupby('CustomerID')['InvoiceNo'].nunique()

plt.figure(figsize=(8, 6))
plt.hist(orders_per_customer, bins=30, color='skyblue', edgecolor='black')
plt.xlabel('Number of Orders')
plt.ylabel('Frequency')
plt.title('Distribution of Orders per Customer')
plt.grid(True)
plt.show()
```



Q. Can you identify the top 5 customers who have made the most purchases by order count?

```
[65]: top_customers = orders_per_customer.sort_values(ascending=False).head(5)
print("Top 5 Customers by Order Count:")
print(top_customers)
```

Top 5 Customers by Order Count:

CustomerID

14911.0 279

12748.0 257

17841.0 208

13089.0 145

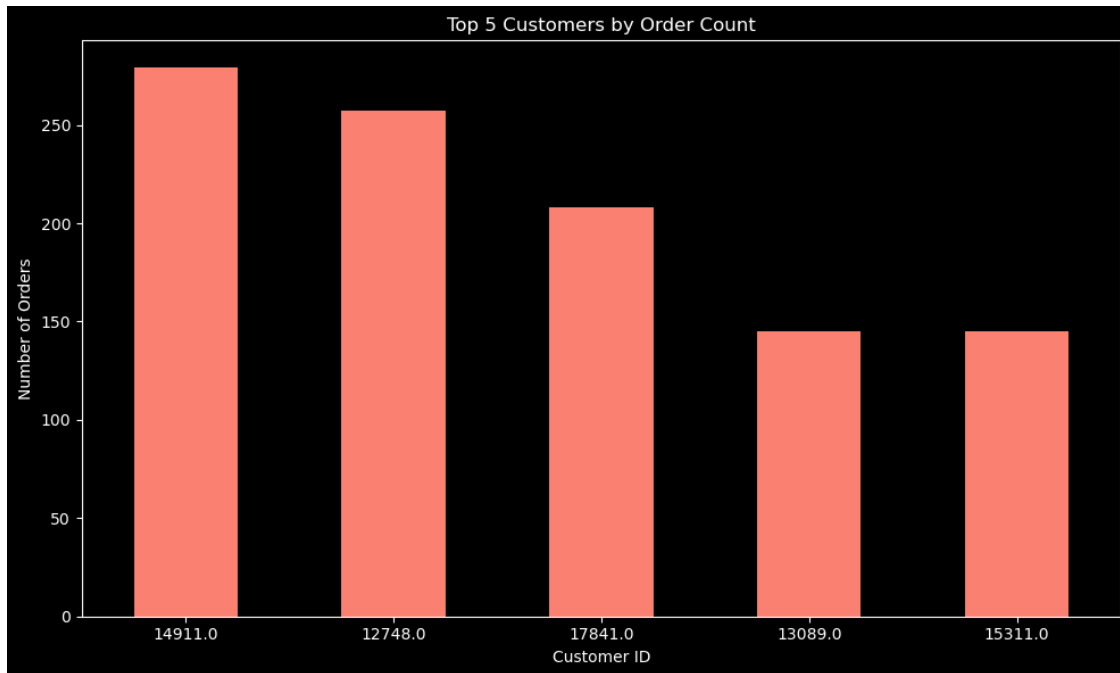
15311.0 145

Name: InvoiceNo, dtype: int64

```
[66]: top_5_customers = orders_per_customer.sort_values(ascending=False).head(5)
plt.figure(figsize=(10, 6))
top_5_customers.plot(kind='bar', color='salmon')
plt.title('Top 5 Customers by Order Count')
```



```
plt.xlabel('Customer ID')
plt.ylabel('Number of Orders')
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()
```



3. PRODUCT ANALYSIS

Q. What are the top 10 most frequently purchased products?

```
[67]: top_10_products = data['Description'].value_counts().head(10)
print("Top 10 Most Frequently Purchased Products:")
print(top_10_products)
```

```
Top 10 Most Frequently Purchased Products:
WHITE HANGING HEART T-LIGHT HOLDER    2369
REGENCY CAKESTAND 3 TIER                2200
JUMBO BAG RED RETROSPOT                 2159
PARTY BUNTING                         1727
LUNCH BAG RED RETROSPOT                 1638
ASSORTED COLOUR BIRD ORNAMENT           1501
SET OF 3 CAKE TINS PANTRY DESIGN         1473
Unknown                                1454
PACK OF 72 RETROSPOT CAKE CASES         1385
LUNCH BAG BLACK SKULL.                  1350
Name: Description, dtype: int64
```

Q. What is the average price of products in the dataset?

```
[68]: average_price = data['UnitPrice'].mean()
      print("Average Price of Products: {:.2f}".format(average_price))
```

Average Price of Products: 4.61

Q. Can you find out which product category generates the highest revenue?

```
[69]: data['TotalRevenue'] = data['Quantity'] * data['UnitPrice']
      highest_revenue_category = data.groupby('Description')['TotalRevenue'].sum().
      ↪idxmax()
      print("\nProduct Category Generating the Highest Revenue: ",
      ↪highest_revenue_category)
```


Product Category Generating the Highest Revenue: DOTCOM POSTAGE

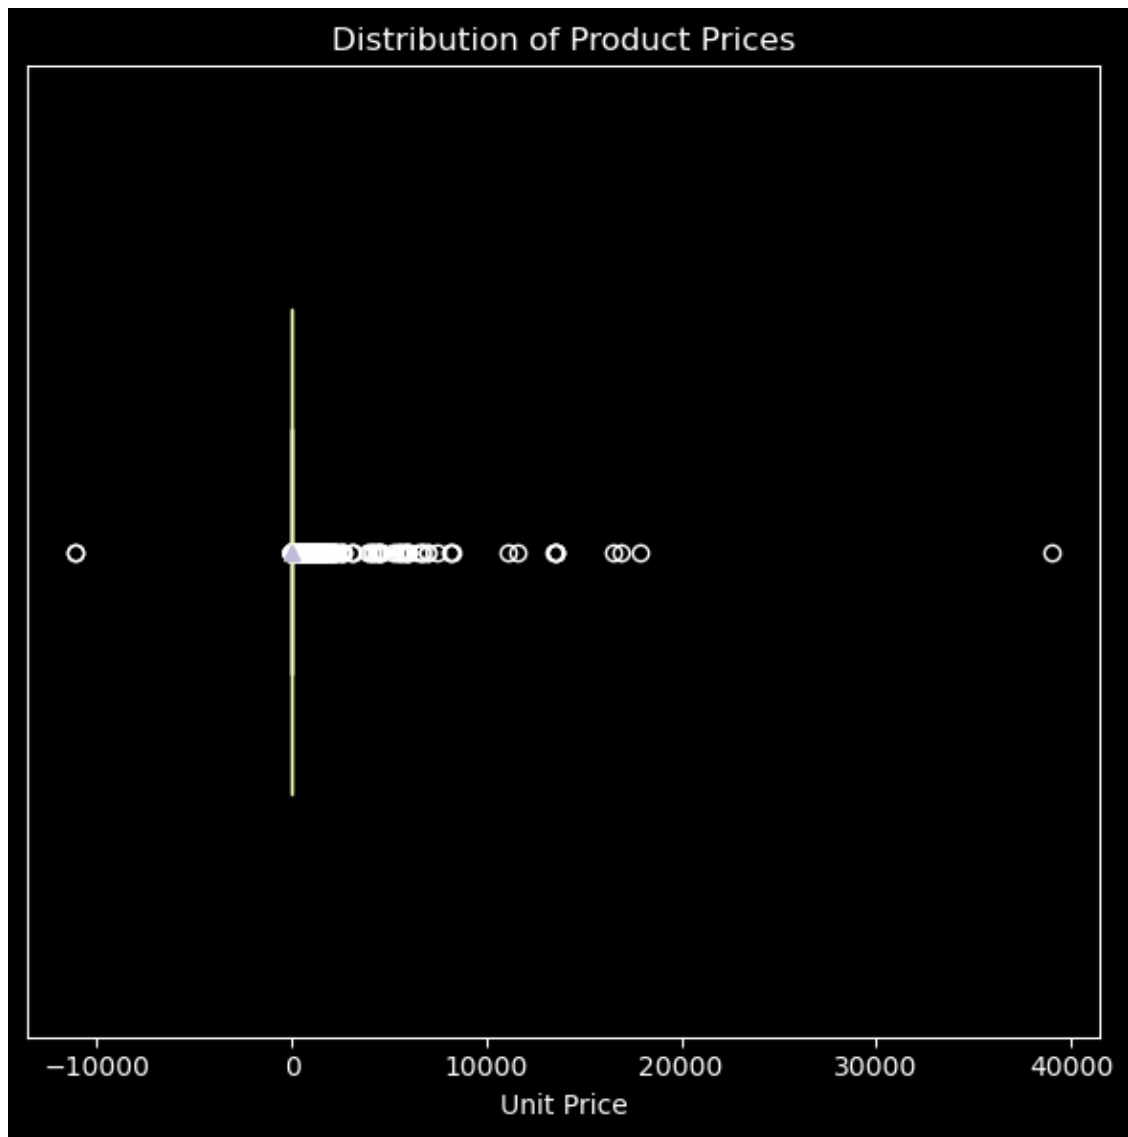
```
[70]: # Top 10 most frequently purchased products
      top_10_products = data['Description'].value_counts().head(10)
      plt.figure(figsize=(10, 6))
      top_10_products.plot(kind='bar', color='skyblue')
      plt.title('Top 10 Most Frequently Purchased Products')
      plt.xlabel('Product')
      plt.ylabel('Frequency')
      plt.xticks(rotation=45)
      plt.tight_layout()
      plt.show()

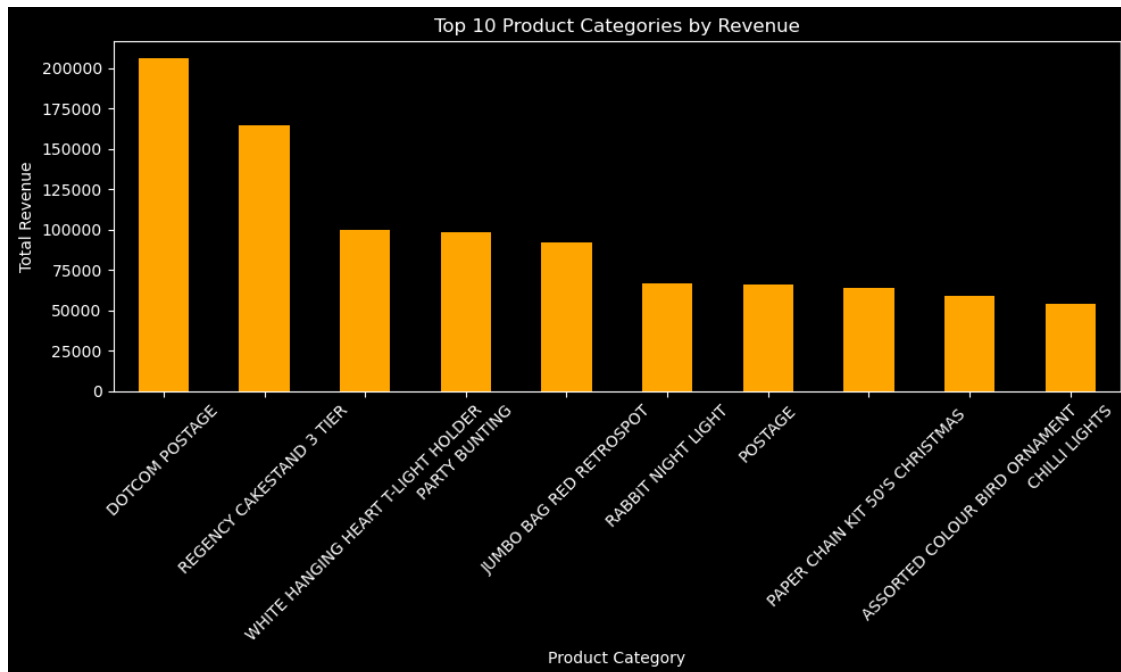
      # Average price of products
      plt.figure(figsize=(6, 6))
      plt.boxplot(data['UnitPrice'], patch_artist=True, showmeans=True, vert=False,
      ↪widths=0.5)
      plt.title('Distribution of Product Prices')
      plt.xlabel('Unit Price')
      plt.yticks([])
      plt.tight_layout()
      plt.show()

      # Product category with highest revenue
      category_revenue = data.groupby('Description')['TotalRevenue'].sum().
      ↪nlargest(10)
      plt.figure(figsize=(10, 6))
      category_revenue.plot(kind='bar', color='orange')
      plt.title('Top 10 Product Categories by Revenue')
      plt.xlabel('Product Category')
      plt.ylabel('Total Revenue')
      plt.xticks(rotation=45)
      plt.tight_layout()
```

```
plt.show()
```







4. TIME ANALYSIS

Q. Is there a specific day of the week or time of day when most orders are placed?

```
[71]: data['InvoiceDate'] = pd.to_datetime(data['InvoiceDate'])

data['DayOfWeek'] = data['InvoiceDate'].dt.day_name()
most_orders_day = data['DayOfWeek'].value_counts().idxmax()
print("Day of the week with the most orders:", most_orders_day)

data['HourOfDay'] = data['InvoiceDate'].dt.hour
most_orders_hour = data['HourOfDay'].value_counts().idxmax()
print("Hour of the day when most orders are placed:", most_orders_hour)
```

Day of the week with the most orders: Thursday

Hour of the day when most orders are placed: 12

Q. What is the average order processing time?

```
[72]: data['OrderProcessingTime'] = data.groupby('InvoiceNo')['InvoiceDate'].
      ↪transform('max') - data.groupby('InvoiceNo')['InvoiceDate'].transform('min')
average_order_processing_time = data['OrderProcessingTime'].mean()
print("Average Order Processing Time:", average_order_processing_time)
```

Average Order Processing Time: 0 days 00:00:00.370578824

Q. Are there any seasonal trends in the dataset?

```
[73]: data['Month'] = data['InvoiceDate'].dt.to_period('M')
monthly_order_count = data['Month'].value_counts().sort_index()

print("\nMonthly Order Counts:")
print(monthly_order_count)
```

Monthly Order Counts:

| | |
|---------|-------|
| 2010-12 | 42481 |
| 2011-01 | 35147 |
| 2011-02 | 27707 |
| 2011-03 | 36748 |
| 2011-04 | 29916 |
| 2011-05 | 37030 |
| 2011-06 | 36874 |
| 2011-07 | 39518 |
| 2011-08 | 35284 |
| 2011-09 | 50226 |
| 2011-10 | 60742 |
| 2011-11 | 84711 |
| 2011-12 | 25525 |


Freq: M, Name: Month, dtype: int64


```
[74]: # Day of the week with the most orders
plt.figure(figsize=(8, 6))
data['DayOfWeek'].value_counts().plot(kind='bar', color='skyblue')
plt.xlabel('Day of the Week')
plt.ylabel('Number of Orders')
plt.title('Orders by Day of the Week')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Hour of the day when most orders are placed
plt.figure(figsize=(8, 6))
data['HourOfDay'].value_counts().sort_index().plot(kind='bar', color='salmon')
plt.xlabel('Hour of the Day')
plt.ylabel('Number of Orders')
plt.title('Orders by Hour of the Day')
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()

# Seasonal Trends
plt.figure(figsize=(10, 6))
monthly_order_count.plot(kind='line', marker='o', color='green')
```

```
plt.xlabel('Month')
plt.ylabel('Number of Orders')
plt.title('Seasonal Trends in Order Counts')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```





5. Geographical Analysis

Q. Can you determine the top 5 countries with the highest number of orders?

```
[75]: top_5_countries_orders = data['Country'].value_counts().head(5)
print("Top 5 Countries by Order Count:")
print(top_5_countries_orders)
```

Top 5 Countries by Order Count:

| | |
|----------------|--------|
| United Kingdom | 495478 |
| Germany | 9495 |
| France | 8557 |
| EIRE | 8196 |
| Spain | 2533 |

Name: Country, dtype: int64

Q. Is there a correlation between the country of the customer and the average order value?

```
[76]: from scipy.stats import f_oneway
avg_order_values = data.groupby('Country')['UnitPrice'].mean()
result_anova = f_oneway(*[data['UnitPrice'][data['Country'] == country] for
    ↪country in data['Country'].unique()])

print("ANOVA Test p-value:", result_anova.pvalue)
```


ANOVA Test p-value: 3.5817635058736703e-47

The p-value obtained from the ANOVA test is extremely low (3.58e-47), shows that it is a highly significant result. The low p-value shows that there are statistically significant differences in average order values among different countries.

Based on the ANOVA test results, we can conclude that there is a correlation between the country of the customer and the average order value. The variation in average order values across various countries is not due to random chance, instead, it suggests that the country of the customer influences the average order value significantly.

This information implies that customers from different countries tend to have varying average spending habits or purchase behaviors, leading to differences in the average value of orders placed by customers from different geographical locations.

```
[77]: import matplotlib.pyplot as plt
plt.figure(figsize=(8, 6))
top_5_countries_orders.plot(kind='bar', color='pink')
plt.xlabel('Country')
plt.ylabel('Number of Orders')
plt.title('Top 5 Countries by Order Count')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```




```
[78]: avg_order_value = data.groupby('Country')['UnitPrice'].mean()

plt.figure(figsize=(12, 6))
bars = plt.bar(avg_order_value.index, avg_order_value.values, color='orange')

plt.xlabel('Country')
plt.ylabel('Average Order Value')
plt.title('Average Order Value by Country')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()

plt.subplots_adjust(bottom=0.4)
for bar in bars:
    plt.text(bar.get_x() + bar.get_width() / 2 - 0.15, bar.get_height() + 0.2,
             round(bar.get_height(), 2), ha='center', va='bottom')

plt.show()
```



6. Payment Analysis

Q. What are the most common payment methods used by customers?

1.2 There is no payment method mentioned in the database. We will be adding a new column for the payment methods, for the sake of analysis.

```
[79]: import random
payment_methods = ['Cash', 'Card', 'Cryptocurrency', 'Mobile Wallet', 'Bank_
↳Transfer']
data['PaymentMethod'] = [random.choice(payment_methods) for _ in_
↳range(len(data))]
```

```
[80]: common_payment_methods = data['PaymentMethod'].value_counts()
print("Most Common Payment Methods:")
print(common_payment_methods)
```

```
Most Common Payment Methods:
Cash                108564
Bank Transfer       108510
Mobile Wallet       108435
Cryptocurrency      108370
Card                108030
Name: PaymentMethod, dtype: int64
```

Q. Is there a relationship between the payment method and the order amount?

```
[81]: payment_order_amount = data.groupby('PaymentMethod')['UnitPrice'].mean()
print("\nMean Order Amount per Payment Method:")
print(payment_order_amount)
```

Mean Order Amount per Payment Method:

PaymentMethod

Bank Transfer 4.439228

Card 5.230973

Cash 4.510574

Cryptocurrency 4.161268

Mobile Wallet 4.715809

Name: UnitPrice, dtype: float64

```
[82]: from scipy.stats import f_oneway
bank_transfer = data[data['PaymentMethod'] == 'Bank Transfer']['UnitPrice']
card = data[data['PaymentMethod'] == 'Card']['UnitPrice']
cash = data[data['PaymentMethod'] == 'Cash']['UnitPrice']
cryptocurrency = data[data['PaymentMethod'] == 'Cryptocurrency']['UnitPrice']
mobile_wallet = data[data['PaymentMethod'] == 'Mobile Wallet']['UnitPrice']

f_statistic, p_value = f_oneway(bank_transfer, card, cash, cryptocurrency,
    ↪mobile_wallet)
print("ANOVA Test p-value:", p_value)
```


ANOVA Test p-value: 0.11793511552775524

A p-value of 0.60 from the ANOVA test suggests that there isn't strong evidence to reject the null hypothesis. In this case, with a higher p-value (greater than the typical significance level of 0.05), it indicates that there may not be significant differences in mean order amounts between the payment methods.


Therefore, based on this analysis, there might not be a statistically significant relationship between the payment method and the order amount.

```
[83]: payment_counts = data['PaymentMethod'].value_counts()

plt.figure(figsize=(8, 6))
payment_counts.plot(kind='bar', color='purple')
plt.xlabel('Payment Method')
plt.ylabel('Frequency')
plt.title('Most Common Payment Methods')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
[84]: plt.figure(figsize=(10, 6))
sns.boxplot(x='PaymentMethod', y='UnitPrice', data=data, palette='coolwarm')
plt.xlabel('Payment Method')
plt.ylabel('Order Amount')
plt.title('Relationship between Payment Method and Order Amount')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



7. CUSTOMER BEHAVIOUR


Q. How long, on average, do customers remain active (between their first and last purchase)?

```
[85]: customer_activity = data.groupby('CustomerID')['InvoiceDate'].agg(['min',
    ↳ 'max'])
customer_activity['ActivityDuration'] = (customer_activity['max'] -
    ↳ customer_activity['min']).dt.days

average_activity_duration = customer_activity['ActivityDuration'].mean()
print("Average time for which the customers remain active:",
    ↳ average_activity_duration, "days")
```

Average time for which the customers remain active: 133.38677950594695 days

```
[86]: plt.figure(figsize=(8, 6))
plt.hist(customer_activity['ActivityDuration'], bins=30, color='skyblue',
    ↳ edgecolor='black')
plt.xlabel('Activity Duration (days)')
plt.ylabel('Frequency')
plt.title('Customer Activity Duration Distribution')
plt.grid(True)
plt.show()
```



8. RETURNS AND REFUNDS

1.2.1 There is no RETURN/REFUND and product category mentioned in the database. We will be adding a new column for the sake of analysis.

```
[87]: data['ReturnRefund'] = np.random.choice(['Return', 'Refund', 'None', 'Null', 'No', ' '], size=len(data))

product_descriptions = [
    'Stylish backpack', 'Antique vase', 'Modern lamp', 'Designer chair', 'Vintage clock', 'Bottle' ]
data['ProductDescription'] = np.random.choice(product_descriptions, size=len(data))

print(data.head())
```

| | InvoiceNo | StockCode | Description | Quantity | \ |
|---|-----------|-----------|------------------------------------|----------|---|
| 0 | 536365 | 85123A | WHITE HANGING HEART T-LIGHT HOLDER | 6 | |
| 1 | 536365 | 71053 | WHITE METAL LANTERN | 6 | |
| 2 | 536365 | 84406B | CREAM CUPID HEARTS COAT HANGER | 8 | |

| | | | | |
|---|--------|--------|-------------------------------------|---|
| 3 | 536365 | 84029G | KNITTED UNION FLAG HOT WATER BOTTLE | 6 |
| 4 | 536365 | 84029E | RED WOOLLY HOTTIE WHITE HEART. | 6 |

| | InvoiceDate | UnitPrice | CustomerID | Country | Date | Time | \ |
|---|---------------------|-----------|------------|----------------|------------|------|---|
| 0 | 2010-12-01 08:26:00 | 2.55 | 17850.0 | United Kingdom | 2010-12-01 | 8:26 | |
| 1 | 2010-12-01 08:26:00 | 3.39 | 17850.0 | United Kingdom | 2010-12-01 | 8:26 | |
| 2 | 2010-12-01 08:26:00 | 2.75 | 17850.0 | United Kingdom | 2010-12-01 | 8:26 | |
| 3 | 2010-12-01 08:26:00 | 3.39 | 17850.0 | United Kingdom | 2010-12-01 | 8:26 | |
| 4 | 2010-12-01 08:26:00 | 3.39 | 17850.0 | United Kingdom | 2010-12-01 | 8:26 | |

| | Recency | TotalPrice | TotalRevenue | DayOfWeek | HourOfDay | \ |
|---|---------|------------|--------------|-----------|-----------|---|
| 0 | 373 | 15.30 | 15.30 | Wednesday | 8 | |
| 1 | 373 | 20.34 | 20.34 | Wednesday | 8 | |
| 2 | 373 | 22.00 | 22.00 | Wednesday | 8 | |
| 3 | 373 | 20.34 | 20.34 | Wednesday | 8 | |
| 4 | 373 | 20.34 | 20.34 | Wednesday | 8 | |

| | OrderProcessingTime | Month | PaymentMethod | ReturnRefund | ProductDescription |
|---|---------------------|---------|----------------|--------------|--------------------|
| 0 | 0 days | 2010-12 | Cash | No | Modern lamp |
| 1 | 0 days | 2010-12 | Mobile Wallet | Return | Vintage clock |
| 2 | 0 days | 2010-12 | Bank Transfer | No | Vintage clock |
| 3 | 0 days | 2010-12 | Bank Transfer | No | Designer chair |
| 4 | 0 days | 2010-12 | Cryptocurrency | Refund | Designer chair |

Q. What is the percentage of orders that have experienced returns or refunds?

```
[88]: total_orders = len(data)
orders_with_returns = len(data[data['ReturnRefund'] == 'Return']) +
    len(data[data['ReturnRefund'] == 'Refund'])
percentage_return_orders = (orders_with_returns / total_orders) * 100

print(f"Percentage of orders with returns or refunds: {percentage_return_orders:
    .2f}%")
```

Percentage of orders with returns or refunds: 33.23%

Q. Is there a correlation between the product category and the likelihood of returns?

```
[89]: return_by_category = data.groupby('ProductDescription')['ReturnRefund'].
    value_counts(normalize=True).unstack().fillna(0)
print(return_by_category[['Return', 'Refund']])
```

| ReturnRefund | Return | Refund |
|--------------------|----------|----------|
| ProductDescription | | |
| Antique vase | 0.165757 | 0.166308 |
| Bottle | 0.163853 | 0.168639 |
| Designer chair | 0.165812 | 0.165690 |
| Modern lamp | 0.166870 | 0.166737 |
| Stylish backpack | 0.163932 | 0.165959 |

Vintage clock 0.167650 0.166510

```
[90]: from scipy.stats import chi2_contingency

chi2, p, _, _ = chi2_contingency(return_by_category[['Return', 'Refund']])
print(f"Chi-square statistic: {chi2}")
print(f"P-value: {p}")
```


Chi-square statistic: 6.83771175307887e-05

P-value: 0.999999999979435

The obtained p-value from the chi-square test is close to 1. This high p-value suggests strong evidence that there is no significant association between the product categories and the likelihood of returns. In other words, based on the data analyzed, it doesn't appear that different product categories have a notable impact on the likelihood of returns.

```
[91]: return_refund_percentage = data['ReturnRefund'].value_counts(normalize=True) * 100

plt.figure(figsize=(6, 4))
return_refund_percentage.plot(kind='bar', color=['skyblue', 'salmon'])
plt.xlabel('Type of Transaction')
plt.ylabel('Percentage')
plt.title('Percentage of Orders with Returns or Refunds')
plt.xticks(rotation=0)
plt.show()
```



9. Profitability Analysis

Q. Can you calculate the total profit generated by the company during the dataset's time period?

```
[92]: data['TotalCost'] = data['Quantity'] * data['UnitPrice']
total_profit = data['TotalCost'].sum()
data['Profit'] = data['TotalCost'] - data['UnitPrice']
print("Total Profit:", total_profit)
```

Total Profit: 9747747.933999998

Q. What are the top 5 products with the highest profit margins?

```
[93]: data['ProfitMargin'] = (data['Profit'] / data['TotalCost']) * 100

profit_per_product = data.groupby('ProductDescription')['Profit'].sum().
    ↪sort_values(ascending=False)

top_5_profitable_products = data.groupby('ProductDescription')['ProfitMargin'].
    ↪mean().nlargest(5)
print("Top 5 products with highest profit margins:")
print(top_5_profitable_products)
```

Top 5 products with highest profit margins:

```
ProductDescription
Vintage clock      58.091086
Designer chair    58.037085
Modern lamp       57.998625
Stylish backpack  57.939480
Bottle            57.938640
Name: ProfitMargin, dtype: float64
```

10. CUSTOMER SATISFACTION

1.2.2 There is no customer feedback mentioned in the database. We will be adding a new column for the sake of analysis.

Q. Is there any data available on customer feedback or ratings for products or services?

```
[95]: additional_comments = [
    "Fantastic product quality!",
    "Disappointed with the late delivery.",
    "Very satisfied with the service.",
    "The product exceeded my expectations.",
    "Unhappy with the customer service response.",
    "Superb packaging, arrived safely.",
```

"The item doesn't match the description.",
"Smooth transaction, thank you!",
"Appreciate the quick resolution to my issue.",
"Impressed with the product's durability.",
"Terrible experience, won't purchase again.",
"Efficient and reliable service.",
"The product is exactly what I needed.",
"Poorly packaged, item arrived damaged.",
"Thrilled with the purchase!",
"Customer support was extremely helpful.",
"The item is of high quality.",
"Average experience, nothing extraordinary.",
"Highly dissatisfied with the purchase.",
"Excellent value for money!",
"Product arrived earlier than expected, great service!",
"Not as described, quite misleading.",
"Overall, satisfied with the purchase.",
"The product quality doesn't justify the price.",
"Had issues with the payment process.",
"Highly recommended, top-notch quality.",
"The customer service team was very responsive.",
"Disappointed with the lack of variety in the selection.",
"Absolutely love the product, will buy again!",
"Received the wrong item, frustrating experience.",
"Great value for money, excellent deal!",
"Seamless checkout process, very convenient.",
"The product is exactly what I was looking for.",
"Poor communication regarding shipping updates.",
"The packaging was secure, item arrived in perfect condition.",
"Average service, nothing exceptional.",
"The product durability is questionable.",
"Issues with the return policy, quite rigid.",
"Fantastic experience, highly impressed!",
"Not happy with the after-sales service.",
"Impressed with the prompt delivery.",
"Expected better quality, somewhat disappointed.",
"Excellent customer support, very helpful.",
"The item was a great addition to my collection.",
"The delivery process needs improvement.",
"Pleased with the overall experience.",
"Product is overpriced for the quality provided.",
"Great customer service, they resolved my query quickly.",
"Item received didn't match the image on the website.",
"The product met my expectations.",

]

```
additional_comments.extend(additional_comments)
```

```
data['CustomerFeedback'] = np.random.choice(additional_comments, size=len(data))
```

Performing sentimental analysis-

```
[97]: pip install textblob
```

```
Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: textblob in
/Users/study/.local/lib/python3.11/site-packages (0.17.1)
Requirement already satisfied: nltk>=3.1 in
/Users/aishwaryabs/anaconda3/lib/python3.11/site-packages (from textblob)
(3.8.1)
Requirement already satisfied: click in
/Users/aishwaryabs/anaconda3/lib/python3.11/site-packages (from
nltk>=3.1->textblob) (8.0.4)
Requirement already satisfied: joblib in
/Users/aishwaryabs/anaconda3/lib/python3.11/site-packages (from
nltk>=3.1->textblob) (1.2.0)
Requirement already satisfied: regex>=2021.8.3 in
/Users/aishwaryabs/anaconda3/lib/python3.11/site-packages (from
nltk>=3.1->textblob) (2022.7.9)
Requirement already satisfied: tqdm in
/Users/aishwaryabs/anaconda3/lib/python3.11/site-packages (from
nltk>=3.1->textblob) (4.65.0)
Note: you may need to restart the kernel to use updated packages.
```

```
[98]: from textblob import TextBlob
feedback = data['CustomerFeedback']
sentiment_scores = feedback.apply(lambda x: TextBlob(str(x)).sentiment.polarity)

positive_sentiment = sentiment_scores[sentiment_scores > 0].count()
negative_sentiment = sentiment_scores[sentiment_scores < 0].count()
neutral_sentiment = sentiment_scores[sentiment_scores == 0].count()


print("Number of positive sentiments:", positive_sentiment)
print("Number of negative sentiments:", negative_sentiment)
print("Number of neutral sentiments:", neutral_sentiment)
```

```
Number of positive sentiments: 303482
Number of negative sentiments: 119429
Number of neutral sentiments: 118998
```

```
[99]: sentiment_labels = ['Positive', 'Negative', 'Neutral']
sentiment_counts = [positive_sentiment, negative_sentiment, neutral_sentiment]
colors = ['lightblue', 'pink', 'lightgreen']

plt.figure(figsize=(8, 6))
```

```
plt.pie(sentiment_counts, labels=sentiment_labels, colors=colors, autopct='%1.
↪1f%%', startangle=140)
plt.title('Sentiment Analysis of Customer Feedback')
plt.axis('equal')
plt.show()
```



[]: